



CHUNGNAM NATIONAL UNIVERSITY



시스템 프로그래밍

강의 11. 시스템 수준 입출력

교재 10장

Nov. 15, 2010

<http://eslab.cnu.ac.kr>

전달 사항

- 헤르미온느와 론 위즐리의 숙제
- 가을을 맞이해서 해야할 일

Introduction(1)

- Input/Output
 - **process of copying data between main memory and external devices such as disk, terminals, and networks**
- **Input** copies data from an I/O device to main memory
- **Output** copies data from memory to a device
- All language run-time systems provide higher level facilities for performing I/O
 - **ANSI C : standard I/O library, e.g. `printf`, `fopen`**
 - **C++ : `<<("put to")`, `>>("get from")`**
- On Unix, these higher level I/O functions are implemented using system-level *Unix I/O functions* provided by kernel

Introduction (2)

- Most of the time, the higher level I/O functions works well and there is no need to use Unix I/O directly.
- Why bother learning about Unix I/O ?
 - **Understanding Unix I/O will help you understand other systems concepts**
 - ▶ e.g. I/O is integrated into the kernel
 - ▶ I/O plays key role in process creation and execution
 - **Sometimes you have no choice but to use Unix I/O**
 - ▶ e.g. no access to file metadata
- Ch.10 is about Unix I/O and standard I/O and you will learn how to use them reliably

Unix 에서의 파일이란

■ A Unix *file* is a sequence of m bytes:

- $B_0, B_1, \dots, B_k, \dots, B_{m-1}$

■ All I/O devices are represented as files:

- `/dev/sda2` (**/usr disk partition**)

- `/dev/tty2` (**terminal**)

- `/dev/cdrom`

■ Even the kernel is represented as a file:

- `/dev/kmem` (**kernel memory image**)

- `/proc` (**kernel data structures**)

Unix 파일의 종류

■ Regular file

- **Binary or text file.**
- **Unix does not know the difference!**

■ Directory file

- **A file that contains the names and locations of other files**

■ Character special and block special files

- **Terminals (character special) and disks (block special)**

■ FIFO (named pipe)

- **A file type used for interprocess communication(?)**

■ Socket

- **A file type used for network communication between processes**

Unix I/O

- The elegant mapping of files to devices allows kernel to export simple interface called Unix I/O.
- Key Unix idea: All input and output is handled in a consistent and uniform way.
- Basic Unix I/O operations (system calls):
 - **Opening and closing files**
 - ▶ `open()` and `close()`
 - **Changing the *current file position* (seek)**
 - ▶ `lseek` (not discussed)
 - **Reading and writing a file**
 - ▶ `read()` and `write()`

파일 열기

- Opening a file informs the kernel that you are getting ready to access that file.

```
int fd;    /* file descriptor */

if ((fd = open("/etc/hosts", O_RDONLY)) < 0) {
    perror("open");
    exit(1);
}
```

- Returns a small identifying integer *file descriptor*
 - **fd == -1 indicates that an error occurred**
- Each process created by a Unix shell begins life with three open files associated with a terminal:
 - **0: STDIN_FILENO, standard input**
 - **1: STDOUT_FILENO, standard output**
 - **2: STDERR_FILENO, standard error**
 - **check <unistd.h> file**

파일 닫기

- Closing a file informs the kernel that you are finished accessing that file.

```
int fd;      /* file descriptor */
int retval; /* return value */

if ((retval = close(fd)) < 0) {
    perror("close");
    exit(1);
}
```

- Closing an already closed file
 - **is an error**
 - **is a recipe for disaster in threaded programs (more on this later)**
- Moral: Always check return codes, even for seemingly benign functions such as `close()`

파일 읽기

- Reading a file copies bytes from the current file position to memory, and then updates file position.

```
char buf[512];
int fd;          /* file descriptor */
int nbytes;      /* number of bytes read */

/* Open file fd ... */
/* Then read up to 512 bytes from file fd */
if ((nbytes = read(fd, buf, sizeof(buf))) < 0) {
    perror("read");
    exit(1);
}
```

- Returns number of bytes read from file `fd` into `buf`
 - `nbytes < 0` indicates that an error occurred.
 - **short counts** (`nbytes < sizeof(buf)`) are possible and are not errors!

파일 쓰기

- Writing a file copies bytes from memory to the current file position, and then updates current file position.

```
char buf[512];
int fd;          /* file descriptor */
int nbytes;      /* number of bytes read */

/* Open the file fd ... */
/* Then write up to 512 bytes from buf to file fd */
if ((nbytes = write(fd, buf, sizeof(buf))) < 0) {
    perror("write");
    exit(1);
}
```

- Returns number of bytes written from `buf` to file `fd`.
 - `nbytes < 0` indicates that an error occurred.
 - **As with reads, short counts are possible and are not errors!**
- Transfers up to 512 bytes from address `buf` to file `fd`

Unix I/O 예제

- Copying standard input to standard output one byte at a time.

```
int main(void)
{
    char c;
    int len;

    while ((len = read(0 /*stdin*/, &c, 1)) == 1) {
        if (write(1 /*stdout*/, &c, 1) != 1) {
            exit(20);
        }
    }
    if (len < 0) {
        printf("read from stdin failed");
        exit(10);
    }
    exit(0);
}
```

- Note the use of error handling wrappers for read and write (Appendix B).

Short Counts 처리하기

- Short counts can occur in these situations:
 - **Encountering (end-of-file) EOF on reads.**
 - **Reading text lines from a terminal.**
 - **Reading and writing network sockets or Unix pipes.**
- Short counts never occur in these situations:
 - **Reading from disk files (except for EOF)**
 - **Writing to disk files.**
- How should you deal with short counts in your code?
 - **Use the RIO (Robust I/O, REB 11.4) package from your textbook's `csapp.c` file (Appendix B).**

파일의 메타데이터

- *Metadata* is data about data, in this case file data.
- Maintained by kernel, accessed by users with the `stat` and `fstat` functions.

```
/* Metadata returned by the stat and fstat functions */
struct stat {
    dev_t      st_dev;      /* device */
    ino_t      st_ino;      /* inode */
    mode_t     st_mode;     /* protection and file type */
    nlink_t    st_nlink;    /* number of hard links */
    uid_t      st_uid;      /* user ID of owner */
    gid_t      st_gid;      /* group ID of owner */
    dev_t      st_rdev;     /* device type (if inode device) */
    off_t      st_size;     /* total size, in bytes */
    unsigned long st_blksize; /* blocksize for filesystem I/O */
    unsigned long st_blocks; /* number of blocks allocated */
    time_t     st_atime;    /* time of last access */
    time_t     st_mtime;    /* time of last modification */
    time_t     st_ctime;    /* time of last change */
};
```

파일 메타데이터 접근하기

```
/* statcheck.c - Querying and manipulating a file's meta data */  
#include "csapp.h"
```

```
int main (int argc, char **argv)  
{  
    struct stat stat;  
    char *type, *readok;  
  
    Stat(argv[1], &stat);  
    if (S_ISREG(stat.st_mode)) /* file type*/  
        type = "regular";  
    else if (S_ISDIR(stat.st_mode))  
        type = "directory";  
    else  
        type = "other";  
    if ((stat.st_mode & S_IRUSR)) /* OK to read?*/  
        readok = "yes";  
    else  
        readok = "no";  
  
    printf("type: %s, read: %s\n", type, readok);  
    exit(0);  
}
```

```
bass> ./statcheck statcheck.c  
type: regular, read: yes  
bass> chmod 000 statcheck.c  
bass> ./statcheck statcheck.c  
type: regular, read: no
```

디렉토리 접근하기

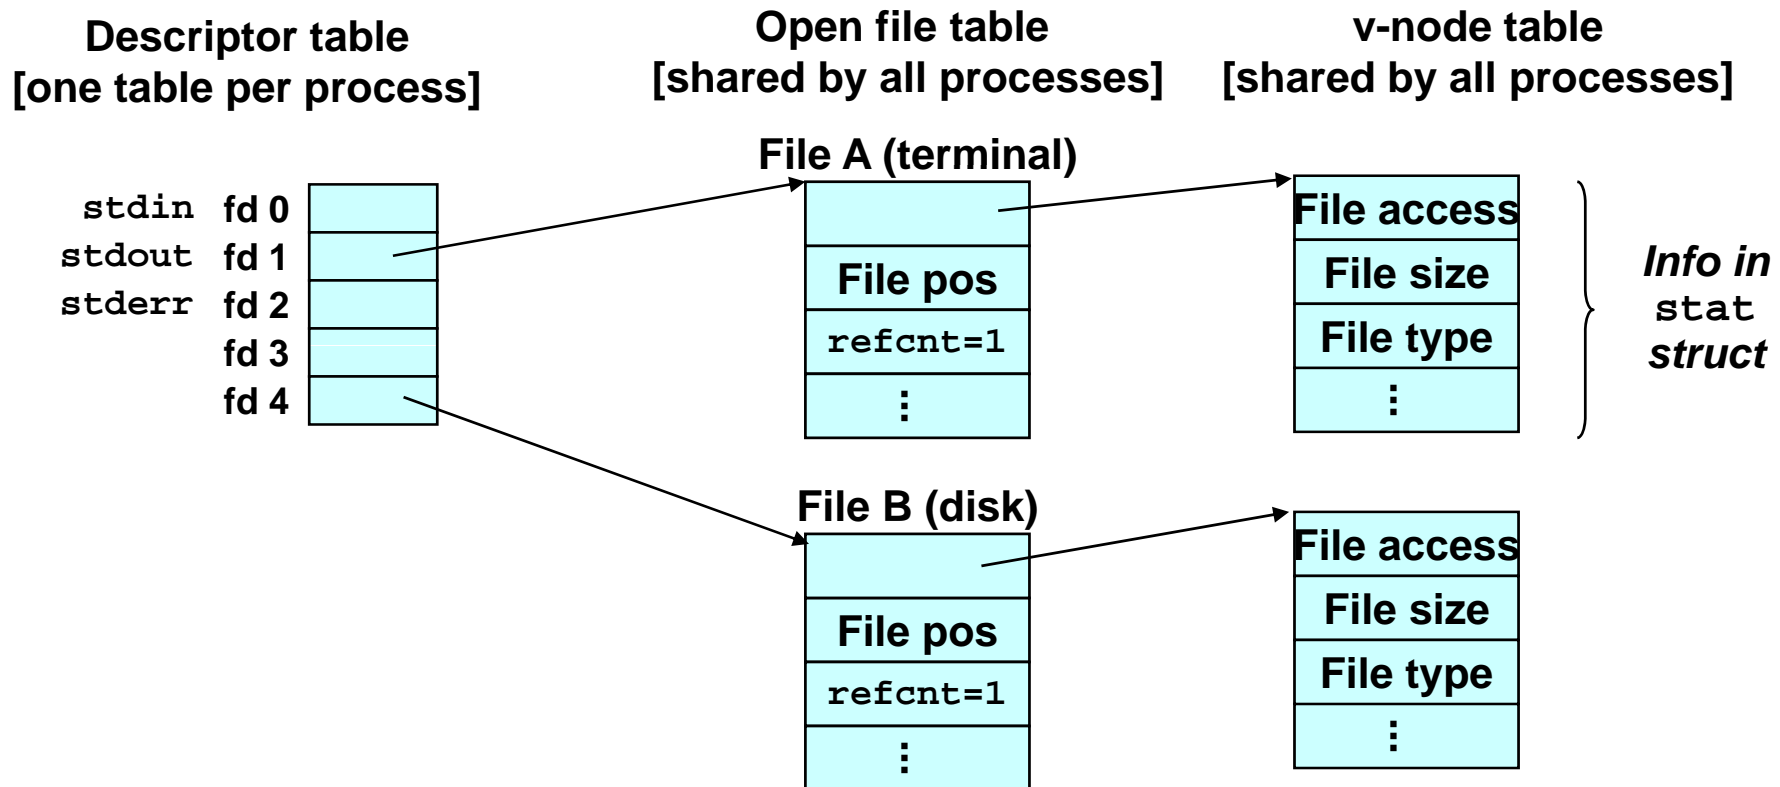
- The only recommended operation on directories is to read its entries.

```
#include <sys/types.h>
#include <dirent.h>

{
    DIR *directory;
    struct dirent *de;
    ...
    if (!(directory = opendir(dir_name)))
        error("Failed to open directory");
    ...
    while (0 != (de = readdir(directory))) {
        printf("Found file: %s\n", de->d_name);
    }
    ...
    closedir(directory);
}
```

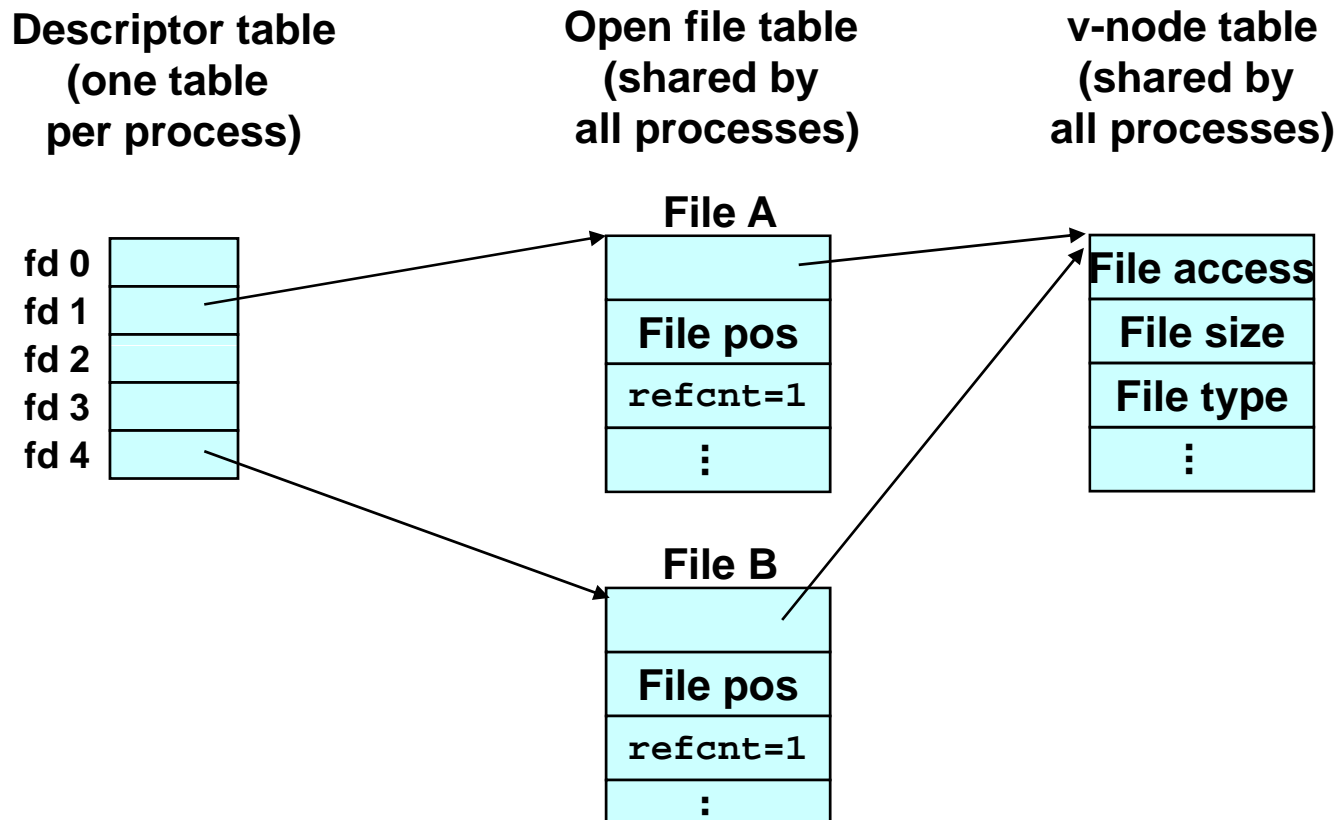

유닉스에서 파일 열기

- Two descriptors referencing two distinct open disk files. Descriptor 1 (stdout) points to terminal, and descriptor 4 points to open disk file.



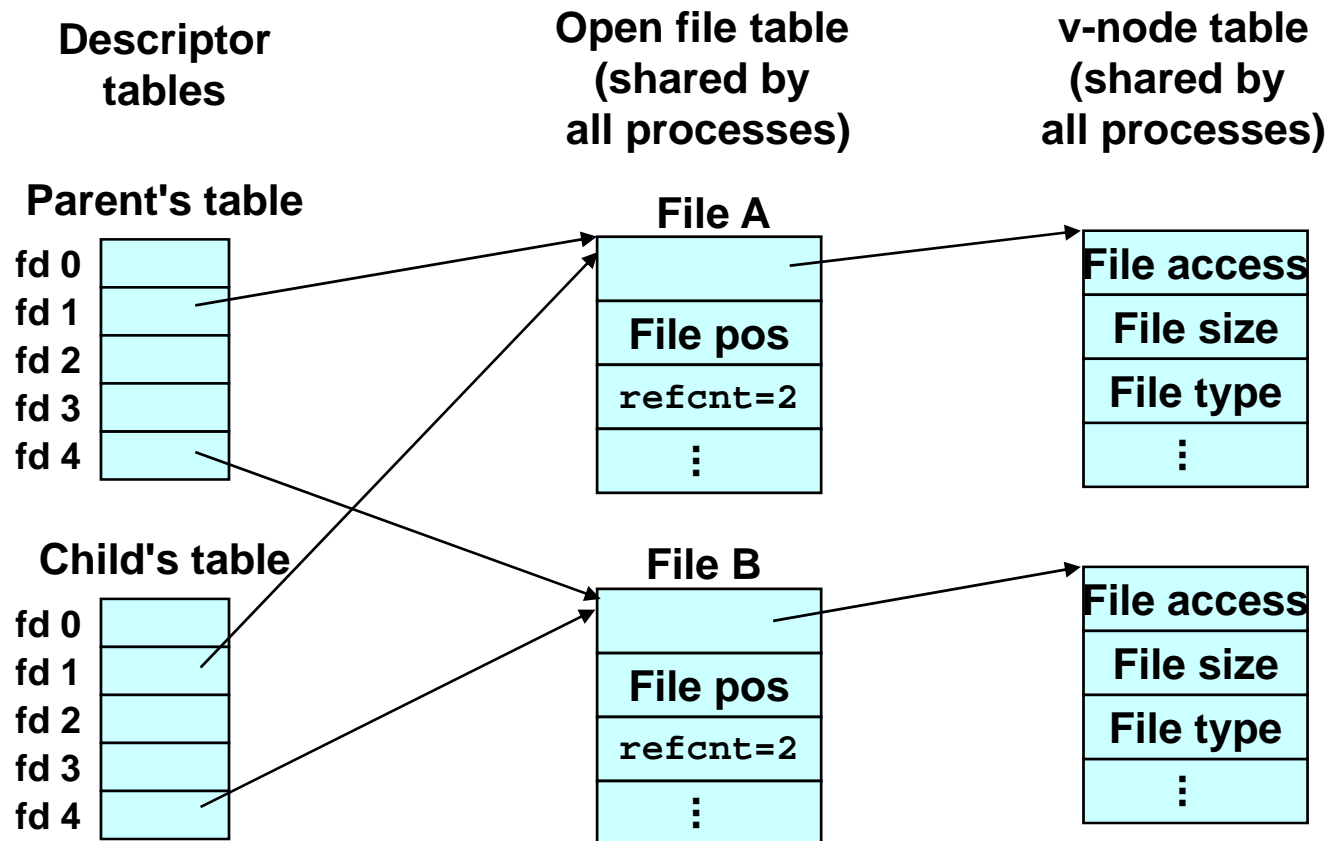
파일의 공유

- Two distinct descriptors sharing the same disk file through two distinct open file table entries
 - E.g., Calling `open` twice with the same `filename` argument



프로세스들 간의 파일의 공유

- A child process inherits its parent's open files. Here is the situation immediately after a `fork`



Problem 11.2

```
int main()
{
    int fd1, fd2;
    char c;

    fd1 = open("foobar.txt", O_RDONLY, 0);
    fd2 = open("foobar.txt", O_RDONLY, 0);
    read(fd1, &c, 1);
    read(fd2, &c, 1);
    printf("c = %c\n", c);
    exit(0);
}
```

foobar.txt contains "foobar", 6 characters.

화면에 무엇이 출력되는가?

Prob. 11.3

```
int main()
{
    int fd;
    char c;

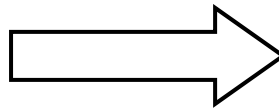
    fd = open("foobar.txt", O_RDONLY, 0);
    if (fork() == 0) {
        read(fd, &c, 1);
        exit(0);
    }
    wait(NULL);
    read(fd, &c, 1);
    printf("c = %c\n", c);
    exit(0);
}
```

입출력에서의 경로변경(Redirection)

- Question: How does a shell implement I/O redirection?
`unix> ls > foo.txt ; redirecting STDOUT to foo.txt`
- Answer: By calling the `dup2(oldfd, newfd)` function
 - Copies (per-process) descriptor table entry `oldfd` to entry `newfd`

Descriptor table
before `dup2(4, 1)`

| | |
|------|---|
| fd 0 | |
| fd 1 | a |
| fd 2 | |
| fd 3 | |
| fd 4 | b |

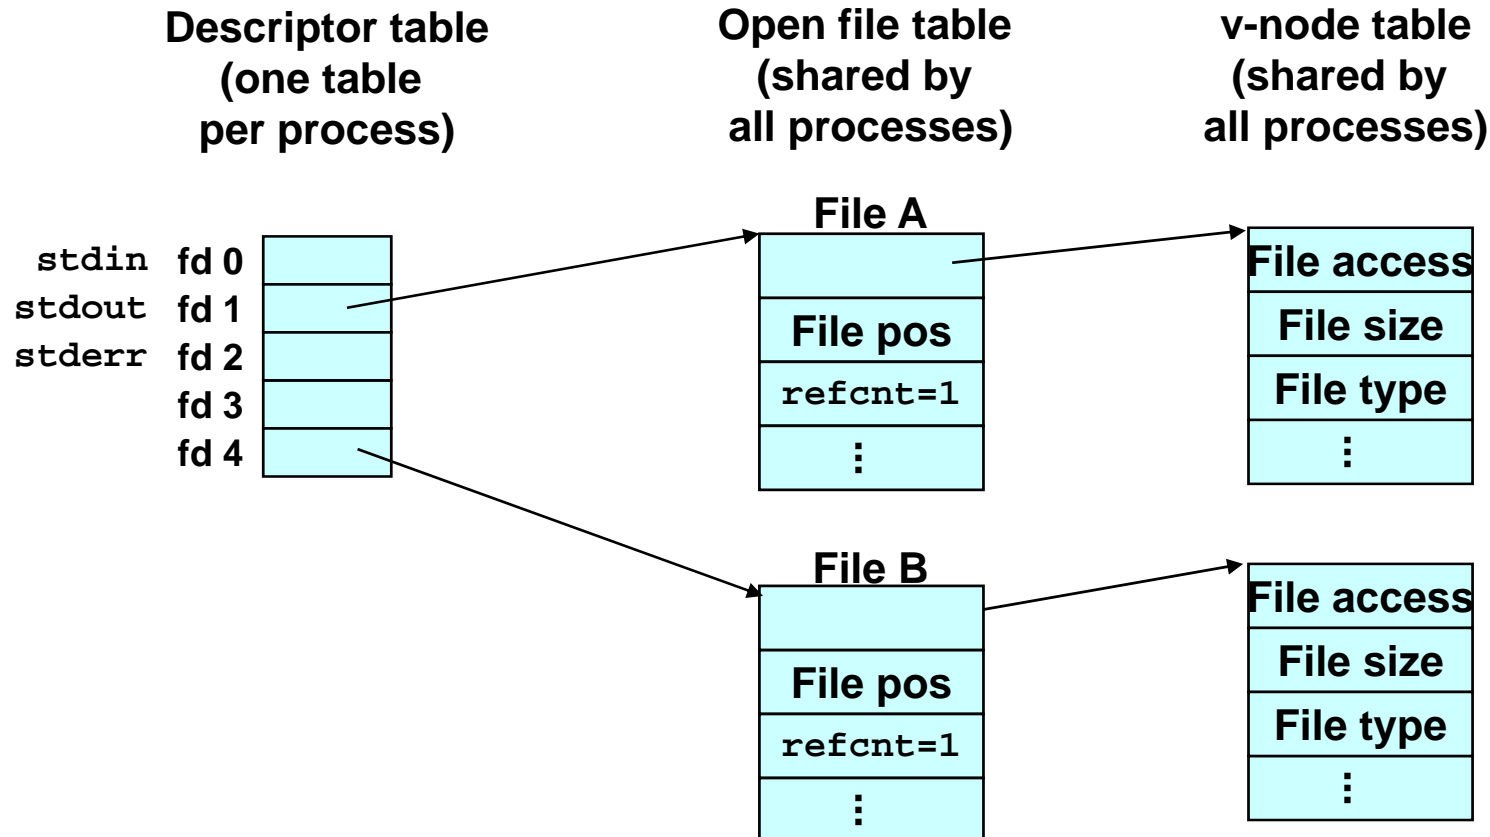


Descriptor table
after `dup2(4, 1)`

| | |
|------|---|
| fd 0 | |
| fd 1 | b |
| fd 2 | |
| fd 3 | |
| fd 4 | b |

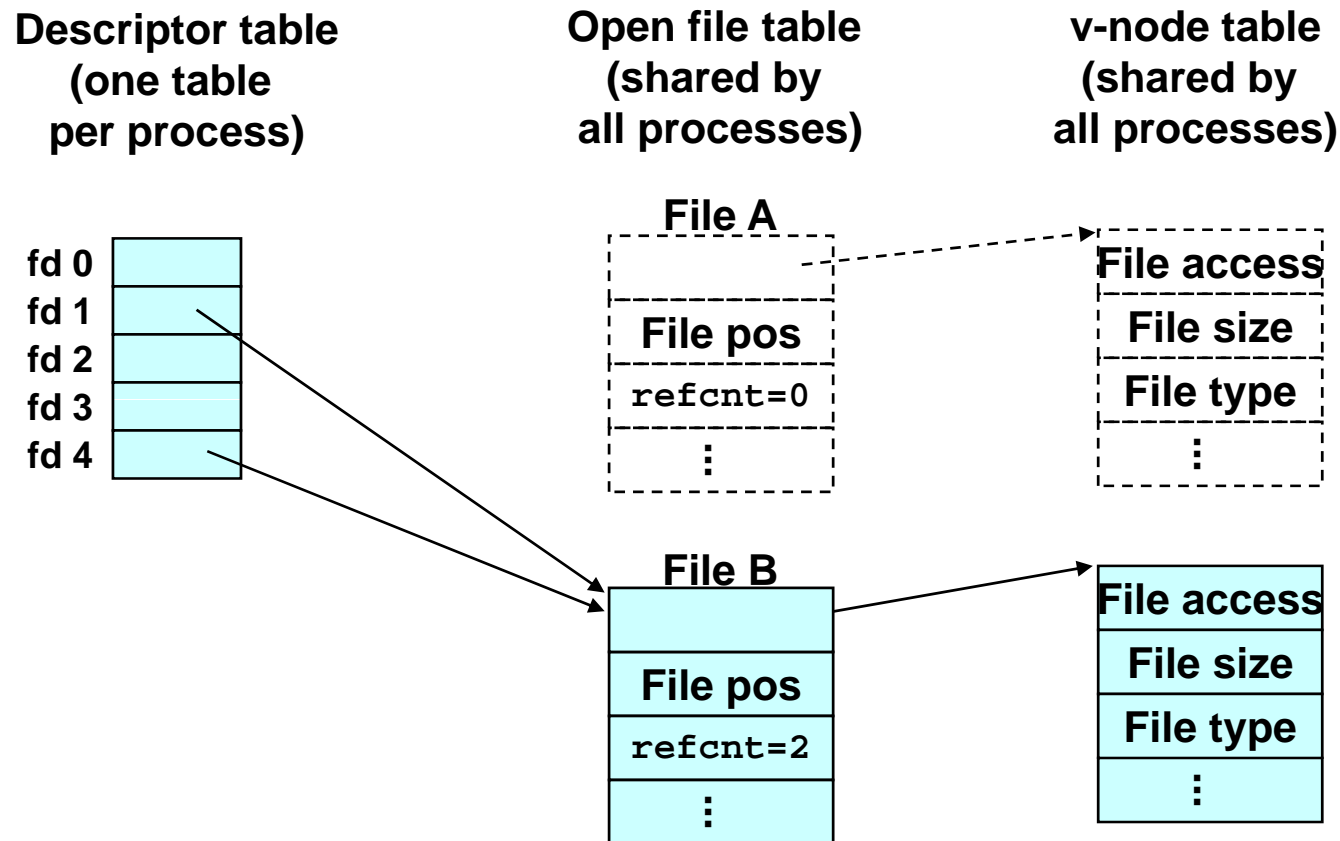
I/O 경로변경 Example

- Before calling `dup2(4, 1)`, `stdout` (descriptor 1) points to a terminal and descriptor 4 points to an open disk file.



I/O 경로변경 Example (cont)

- After calling `dup2(4, 1)`, `stdout` is now redirected to the disk file pointed at by descriptor 4.



표준 입출력 함수

- The C standard library (`libc.a`) contains a collection of higher-level **standard I/O** functions
 - **Documented in Appendix B of K&R.**
- Examples of standard I/O functions:
 - **Opening and closing files (`fopen` and `fclose`)**
 - **Reading and writing bytes (`fread` and `fwrite`)**
 - **Reading and writing text lines (`fgets` and `fputs`)**
 - **Formatted reading and writing (`fscanf` and `fprintf`)**

표준 입출력 스트림

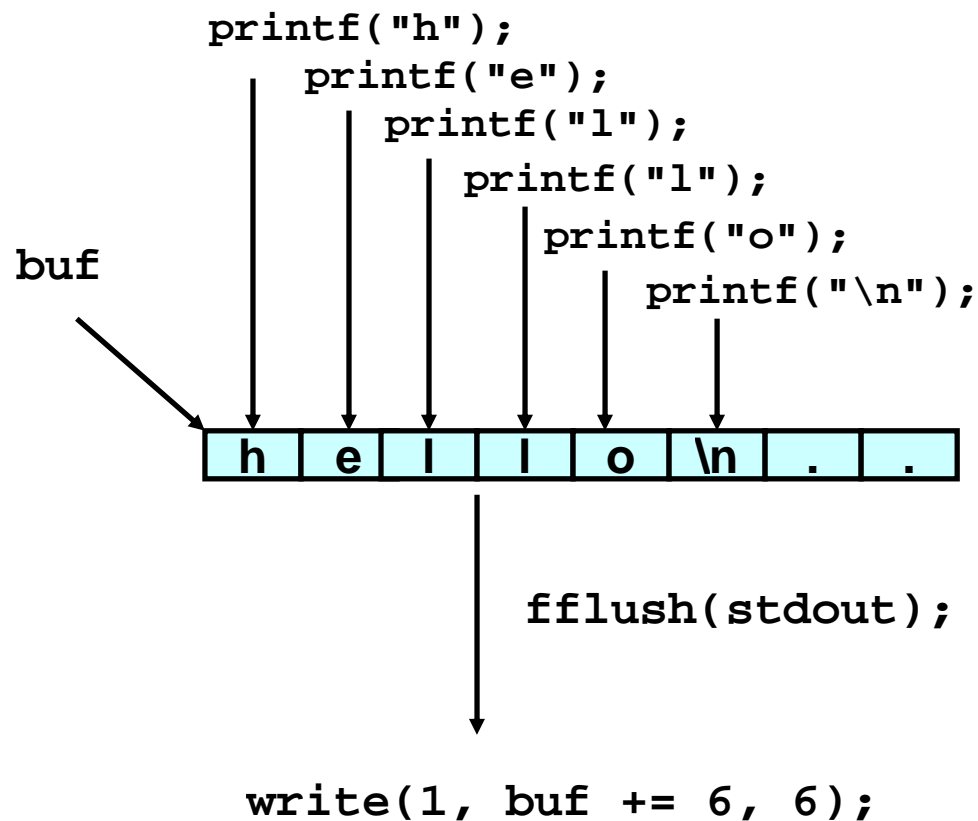
- Standard I/O models open files as *streams*
 - **Abstraction for a file descriptor and a buffer in memory.**
- C programs begin life with three open streams (defined in `stdio.h`)
 - **`stdin` (standard input)**
 - **`stdout` (standard output)**
 - **`stderr` (standard error)**

```
#include <stdio.h>
extern FILE *stdin; /* standard input (descriptor 0) */
extern FILE *stdout; /* standard output (descriptor 1) */
extern FILE *stderr; /* standard error (descriptor 2) */

int main() {
    fprintf(stdout, "Hello, world\n");
}
```

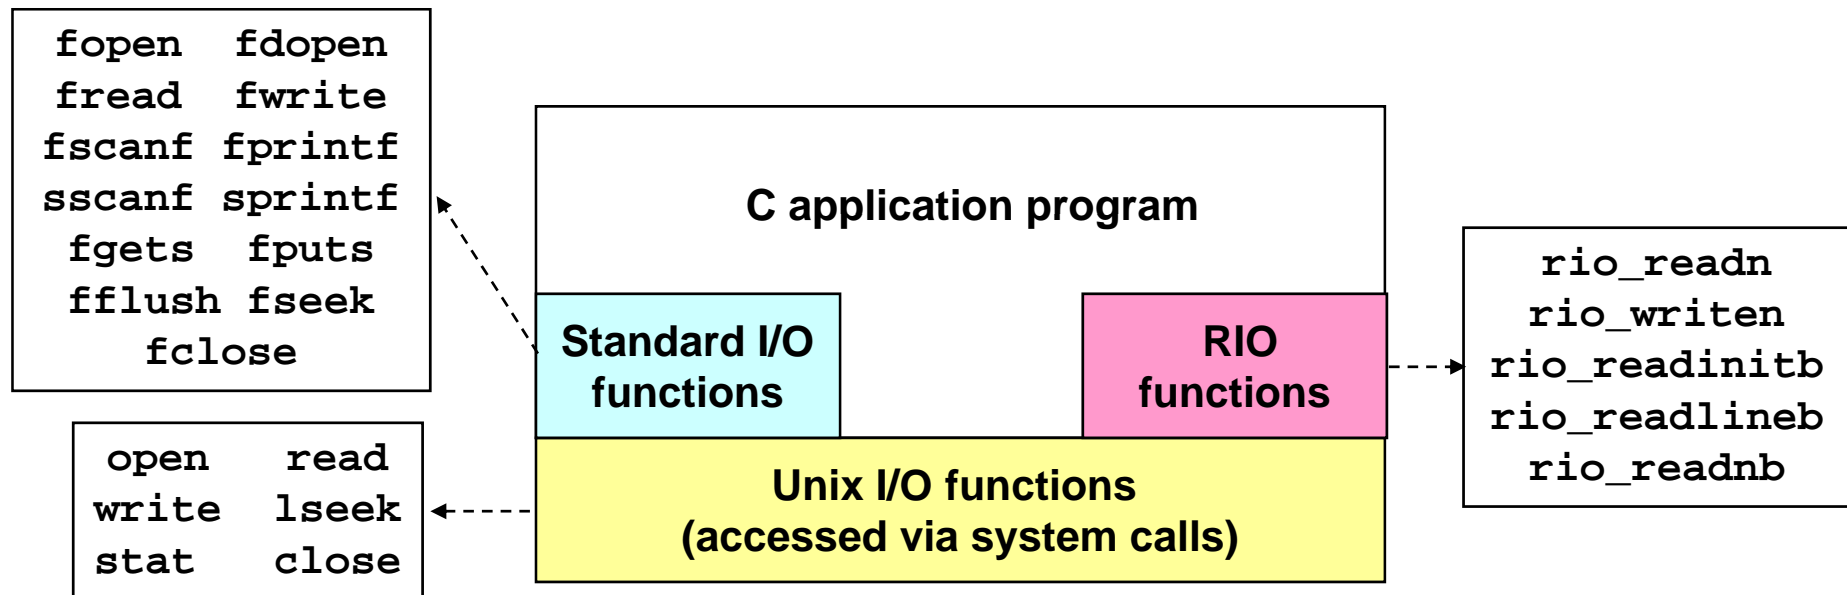
표준 입출력에서의 버퍼링

- Standard I/O functions use buffered I/O



Unix I/O vs. Standard I/O vs. RIO

- Standard I/O and RIO are implemented using low-level Unix I/O.



- Which ones should you use in your programs?

Unix I/O의 장단점

■ Pros

- **Unix I/O is the most general and lowest overhead form of I/O.**
 - ▶ All other I/O packages are implemented using Unix I/O functions.
- **Unix I/O provides functions for accessing file metadata.**

■ Cons

- **Dealing with short counts is tricky and error prone.**
- **Efficient reading of text lines requires some form of buffering, also tricky and error prone.**
- **Both of these issues are addressed by the standard I/O and RIO packages.**

Standard I/O의 장단점

■ Pros:

- **Buffering increases efficiency by decreasing the number of `read` and `write` system calls.**
- **Short counts are handled automatically.**

■ Cons:

- **Provides no function for accessing file metadata**
- **Standard I/O is not appropriate for input and output on network sockets**
- **There are poorly documented restrictions on streams that interact badly with restrictions on sockets**

어느 입출력 함수를 사용할 것인가

- General rule: Use the highest-level I/O functions you can.
 - **Many C programmers are able to do all of their work using the standard I/O functions.**
- When to use standard I/O?
 - **When working with disk or terminal files.**
- When to use raw Unix I/O
 - **When you need to fetch file metadata.**
 - **In rare cases when you need absolute highest performance.**
- When to use RIO?
 - **When you are reading and writing network sockets or pipes.**
 - **Never use standard I/O or raw Unix I/O on sockets or pipes.**



CHUNGNAM NATIONAL UNIVERSITY



프로세스간의 통신 Inter-process Communication

Unix Pipes

* 참고자료 : 관심있는 사람들만 !!!

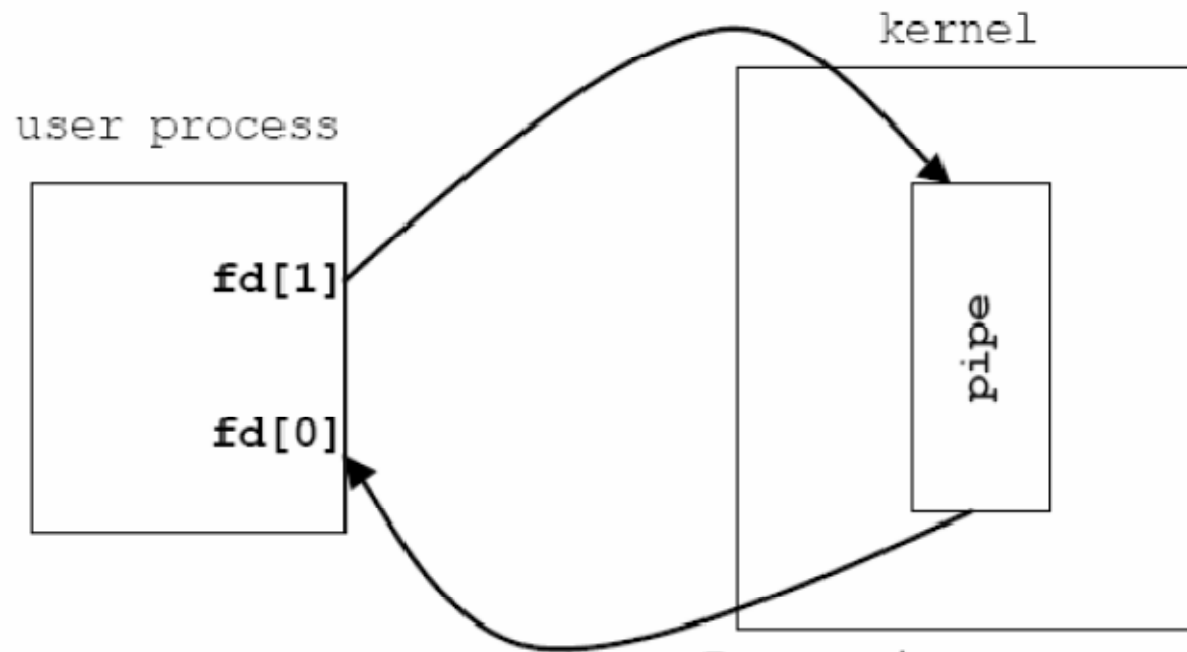
파이프(Pipes)

- The oldest form of Unix IPC and provided by all Unix systems
- Two limitations
 - **Half-duplex : data flows only in one direction**
 - **Can be used only between processes that have a common ancestor**
 - ▶ Usu used between parent and child processes

파이프 생성하기(1)

■ `int pipe (int fd[2])`

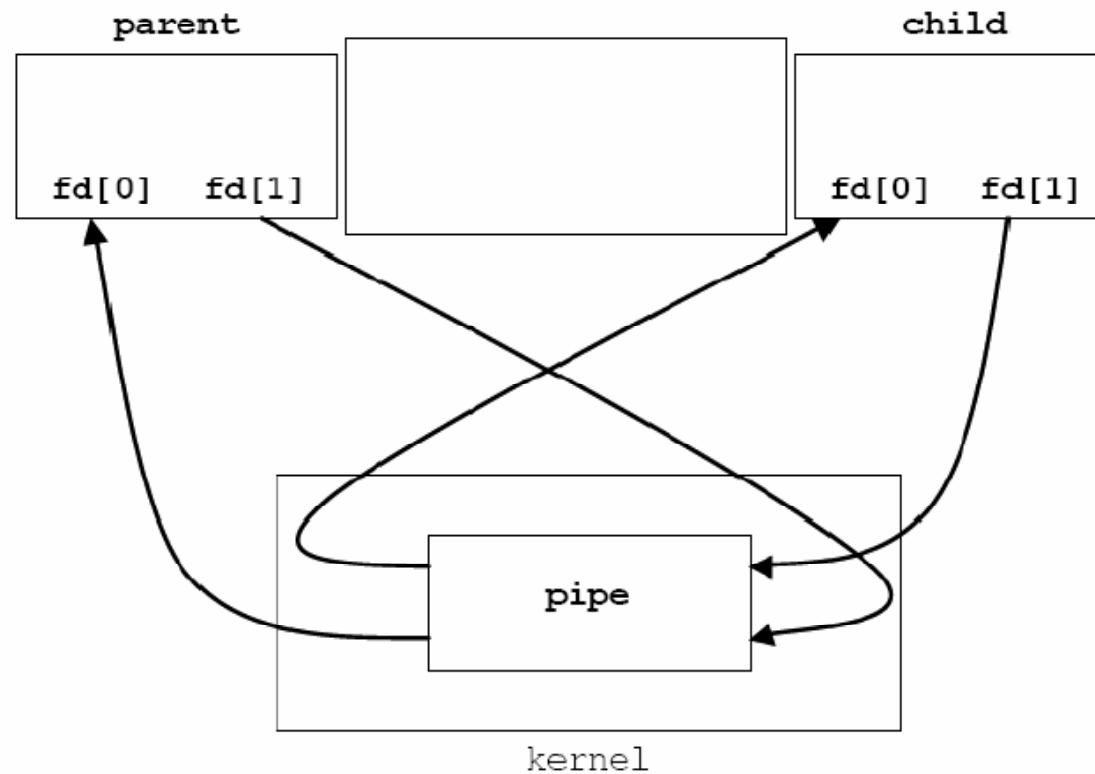
- Two file descriptors are returned through the **fd** argument
 - ▶ `fd[0]`: open for reading
 - ▶ `fd[1]`: open for writing
- The output of **fd[1]** is the input for **fd[0]**.



파이프 생성하기(2)

```
parent => child:  
parent closes fd[0];  
child closes fd[1];
```

```
parent <= child:  
parent closes fd[1];  
child closes fd[0];
```



파이프 읽기/쓰기

- When one end of a pipe is closed,
 - reading from a pipe returns an end of file.
 - writing to a pipe causes **SIGPIPE** is generated and the write returns an error (**EPIPE**).
 - A write of **PIPE_BUF** (kernel's pipe buffer size) bytes or less will not be interleaved with the writes from other processes.
 - **fstat** function returns a file type of FIFO for the pipe file descriptors (can be tested by **S_ISFIFO** macro)

- You should close unused file descriptors!

파이프 사용하기

```
#include <unistd.h>

#define MAXLINE 80

int main (void)
{
    int n, fd[2];
    pid_t pid;
    char line[MAXLINE];

    if (pipe(fd) < 0) exit (1);
    if ((pid = fork()) < 0) exit (2);
    else if (pid > 0) {                /* parent */
        close (fd[0]);
        write (fd[1], "hello world\n", 12);
    }
    else {                            /* child */
        close (fd[1]);
        n = read(fd[0], line, MAXLINE);
        write (1, line, n);
    }
}
```

FIFOs

- `int mkfifo (const char *path, mode_t mode)`
 - Named pipes
 - Unrelated processes can exchange data, whereas pipes can be used only between related processes.
 - FIFO is a type of file: FIFO type (**S_ISFIFO** macro)
 - Once a FIFO created, the normal file I/O functions all work with FIFO.

- `/usr/bin/mkfifo` program can be used to make FIFOs on the command line.

FIFOs 사용하기

■ Opening a FIFO

- An open for read(write)-only blocks until some other process opens the FIFO for writing(reading).

■ Reading/Writing a FIFO

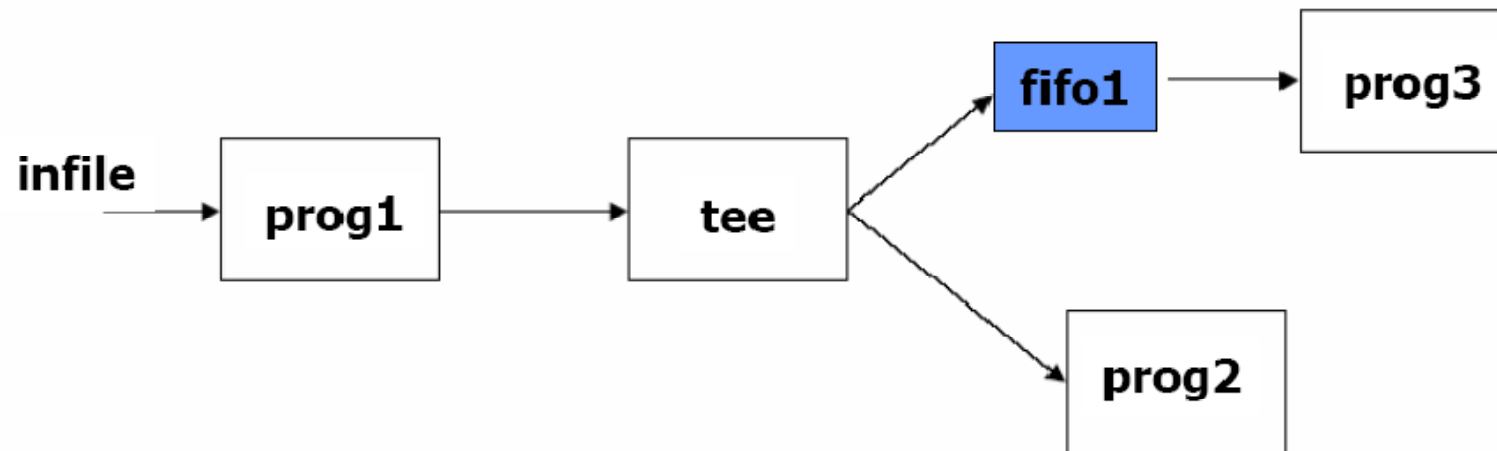
- Writing to a FIFO that no process has open for reading causes **SIGPIPE** to generate.
- When the last writer for a FIFO closes the FIFO, an end of file is generated for the reader of the FIFO.
- **PIPE_BUF**: the maximum amount of data that can be written atomically to a FIFO (without being interleaved among multiple writers).

FIFO의 용도(1)

■ Duplicating a Stream

- Shell commands to pass data from one shell pipeline to another without creating intermediate temporary files

```
$ mkfifo fifo1  
$ prog3 < fifo1 &  
$ prog1 < infile | tee fifo1 | prog2
```



FIFO의 용도(2)

Client-server Communication

- A client-server application to pass data between the clients and server on the same machine.
 - ▶ Clients write to a “well-known” FIFO to send a request to the server.

