



# 시스템 프로그래밍

강의 12. 메모리

교재 9 장

Nov. 19, 2010

<http://eslab.cnu.ac.kr>

# 전달 사항

---

■ 다음주 화요일 수업 휴강 => 오늘 보강했으니

# 학습 내용

- 제 1주 : 과목 소개 및 시스템 개요
- 제 2주 ~ 3주 : 컴퓨터에서의 데이터 표현
  - 정수
  - 실수
- 제 4주~7주 : 어셈블리어와 프로세서 구조
  - 데이터의 이동
  - 제어문
  - 프로시저 및 스택프레임
  - 버퍼 오버플로우
- 제9주 ~ 10주 예외적 인 제어흐름
  - 프로세스
  - 시그널
- 제 11주 입출력 시스템
- 제 12 ~ 13주 메모리 관리
  - 가상메모리
  - 동적 메모리 할당
- 제 14주 동시성 프로그래밍 (Concurrent programming)

# 메모리에 관한 불편한 진실

---


## *Memory Matters*

- Memory is not unbounded
  - **It must be allocated and managed**
  - **Many applications are memory dominated**
    - ▶ Especially those based on complex, graph algorithms
- Memory referencing bugs especially pernicious
  - **Effects are distant in both time and space**
- Memory performance is not uniform
  - **Cache and virtual memory effects can greatly affect program performance**
  - **Adapting program to characteristics of memory system can lead to major speed improvements**

## 동적 메모리 할당 사용의 이유

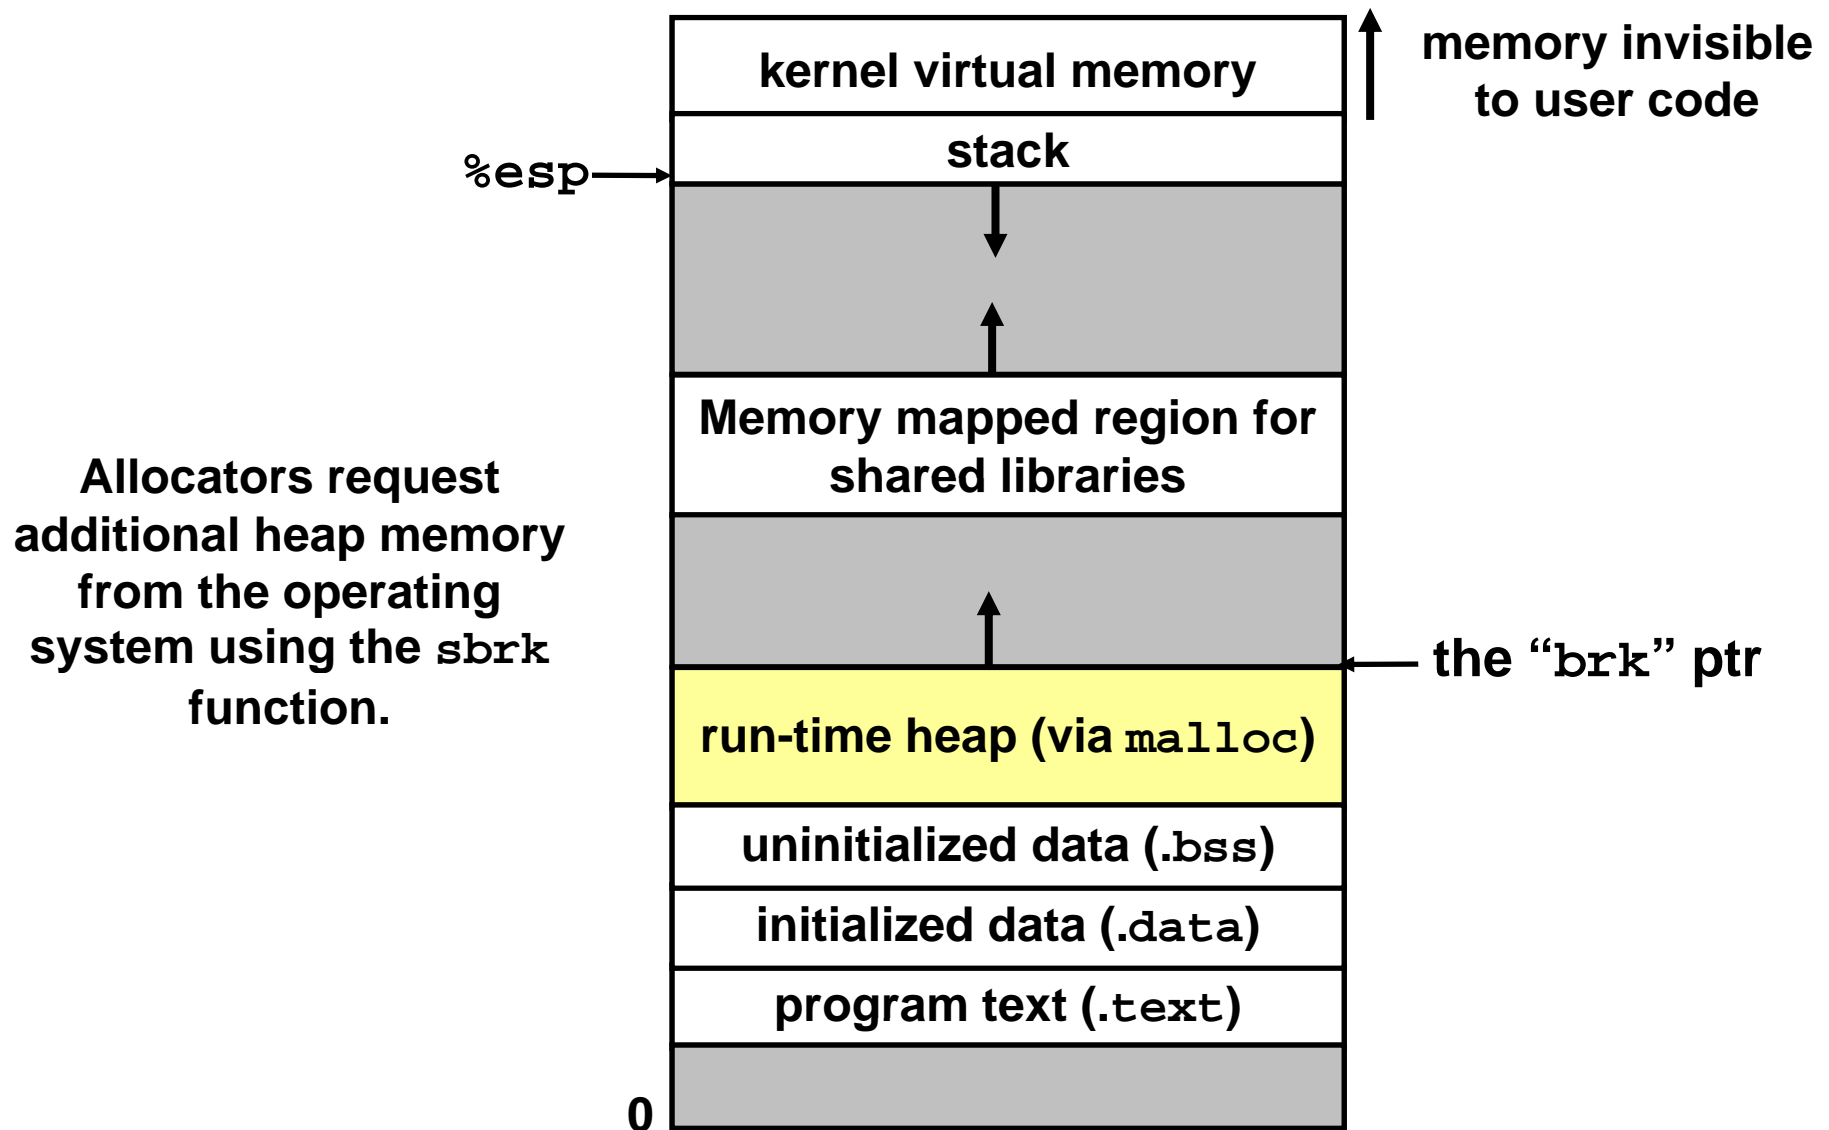
- 프로그램이 실행되기 전에는 크기를 알 수 없는 자료 구조를 위해 사용
- 예 :  $n$  개의 문자를 화면에서 읽어 들여서 배열에 저장하고자 할 때,  $n$  과 문자를 차례로 받아 들여서 실행하는 경우
  - 고정 크기의 배열로도 구현 가능
  - **int array[MAX\_SIZE];**
  - 이와 같이 배열의 크기를 알 수 없을 때, 최대값으로 배열을 구현하는 것은 나쁜 생각
    - ▶ 시스템의 가능한 메모리 사용량을 알 수 없다
    - ▶ MAX\_SIZE 보다 많은 입력을 원한다면?
    - ▶ MAX\_SIZE 값을 계속 바꿔서 다시 컴파일 해야 한다
  - 코드 관리 차원에서 안 좋음
  - 이런 경우에 동적 메모리 할당이 효과적이고, 중요한 프로그래밍 기술임

## 동적 메모리 할당



- 직접 vs 간접(Explicit vs. Implicit) Memory Allocator
  - 직접 할당 : **application allocates and frees space**
    - ▶ E.g., `malloc` and `free` in C
  - 간접 할당 : **application allocates, but does not free space**
    - ▶ E.g. garbage collection in Java, ML or Lisp
- Allocation
  - **In both cases the memory allocator provides an abstraction of memory as a set of blocks**
  - 응용프로그램에 **free** 메모리 블록을 나눠준다
- Will discuss simple explicit memory allocation

## 프로세스의 메모리 이미지



# Malloc 패키지

---

```
#include <stdlib.h>
```

```
void *malloc(size_t size)
```

- **If successful:**

- ▶ Returns a pointer to a memory block of at least `size` bytes, (typically) aligned to 8-byte boundary.
- ▶ If `size == 0`, returns NULL

- **If unsuccessful: returns NULL (0) and sets `errno`.**

```
void free(void *p)
```

- **Returns the block pointed at by `p` to pool of available memory**

- **`p` must come from a previous call to `malloc` or `realloc`.**

```
void *realloc(void *p, size_t size)
```

- **Changes size of block `p` and returns pointer to new block.**

- **Contents of new block unchanged up to min of old and new size.**



# Malloc Example

```
void foo(int n, int m) {
    int i, *p;

    /* allocate a block of n ints */
    p = (int *)malloc(n * sizeof(int));
    if (p == NULL) {
        perror("malloc");
        exit(0);
    }
    for (i=0; i<n; i++) p[i] = i;

    /* add m bytes to end of p block */
    if ((p = (int *) realloc(p, (n+m) * sizeof(int))) == NULL) {
        perror("realloc");
        exit(0);
    }
    for (i=n; i < n+m; i++) p[i] = i;


    /* print new array */
    for (i=0; i<n+m; i++)
        printf("%d\n", p[i]);

    free(p); /* return p to available memory pool */
}
```

# 가정

## ■ Assumptions made in this lecture

- **Memory is word addressed (each word can hold a pointer)**




## 할당 예제


```
p1 = malloc(4)
```




```
p2 = malloc(5)
```




```
p3 = malloc(6)
```



```
free(p2)
```



```
p4 = malloc(2)
```



# 제한사항

## ■ 응용 프로그램

- **Can issue arbitrary sequence of allocation and free requests**
- **Free requests must correspond to an allocated block**

## ■ 할당 프로그램

- **Can't control number or size of allocated blocks**
- **Must respond immediately to all allocation requests**
  - ▶ *i.e.*, can't reorder or buffer requests
- **Must allocate blocks from free memory**
  - ▶ *i.e.*, can only place allocated blocks in free memory
- **Must align blocks so they satisfy all alignment requirements**
  - ▶ 8 byte alignment for GNU malloc (`libc` malloc) on Linux boxes
- **Can only manipulate and modify free memory**
- **Can't move the allocated blocks once they are allocated**
  - ▶ *i.e.*, compaction is not allowed

# 우수한 `malloc/free` 프로그램의 목표

---

## ■ Primary goals

### ● Good time performance for `malloc` and `free`

- ▶ Ideally should take constant time (not always possible)
- ▶ Should certainly not take linear time in the number of blocks

### ● Good space utilization

- ▶ User allocated structures should be large fraction of the heap.
- ▶ Want to minimize “fragmentation”.

## ■ Some other goals

### ● Good locality properties

- ▶ Structures allocated close in time should be close in space
- ▶ “Similar” objects should be allocated close in space

### ● Robust

- ▶ Can check that `free(p1)` is on a valid allocated object `p1`
- ▶ Can check that memory references are to allocated space

## 성능 지표 : Throughput


- Given some sequence of `malloc` and `free` requests:
  - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- Want to maximize throughput and peak memory utilization.
  - **These goals are often conflicting**
- Throughput:
  - **Number of completed requests per unit time**
  - **Example:**
    - ▶ 5,000 `malloc` calls and 5,000 `free` calls in 10 seconds
    - ▶ Throughput is 1,000 operations/second.
  - **`alloc`과 `free` 함수의 평균 처리 시간을 최소화하면 **throughput**을 극대화할 수 있다**

## 성능지표 : 순간 최대 메모리 사용율

- 좋은 프로그래머는 가상메모리의 크기도 제한되어 있다는 것을 알고 작업한다
  - 따라서 효율적으로 관리해야 한다
- 효율적인 heap 사용 지표 : peak heap utilization
- Given some sequence of `malloc` and `free` requests:
  - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- Def: Aggregate payload  $P_k$ :
  - `malloc(p)` results in a block with a **payload** of  $p$  bytes..
  - After request  $R_k$  has completed, the **aggregate payload**  $P_k$  is the sum of currently allocated payloads.
- Def: Current heap size is denoted by  $H_k$ 
  - Assume that  $H_k$  is monotonically nondecreasing
- Def: Peak memory utilization:
  - After  $k$  requests, **peak memory utilization** is:
    - ▶  $U_k = (\max_{i \leq k} P_i) / H_k$
- 구현 목표 :  $U_{n-1}$  을 가능한 모든 작업 순서에 대해서 극대화한다

# 내부 메모리 단편화(Internal Fragmentation)

- Poor memory utilization caused by *fragmentation*.
  - Comes in two forms: internal and external fragmentation
- Internal fragmentation
  - For some block, internal fragmentation is the difference between the block size and the payload size.




- Caused by overhead of maintaining heap data structures, padding for alignment purposes, or explicit policy decisions (e.g., not to split the block).
- Depends only on the pattern of *previous* requests, and thus is easy to measure.




# 외부 메모리 단편화(External Fragmentation)

Occurs when there is enough aggregate heap memory, but no single free block is large enough


```
p1 = malloc(4)
```




```
p2 = malloc(5)
```



```
p3 = malloc(6)
```



```
free(p2)
```




```
p4 = malloc(6)
```

**oops!**

External fragmentation depends on the pattern of *future* requests, and thus is difficult to measure.

## 구현시 고려할 점


- How do we know how much memory to free just given a pointer?
- How do we keep track of the free blocks?
- What do we do with the extra space when allocating a structure that is smaller than the free block it is placed in?
- How do we pick a block to use for allocation -- many might fit?
- How do we reinsert freed block?



# 얼마만 큼 Free 시켜야 하는지 결정하기


## ■ Standard method

- **Keep the length of a block in the word preceding the block.**
  - ▶ This word is often called the *header field* or *header*
- **Requires an extra word for every allocated block**




# Free 블록 관리하기

- Method 1: *Implicit list* using lengths -- links all blocks



- Method 2: *Explicit list* among the free blocks using pointers within the free blocks



- Method 3: *Segregated free list*

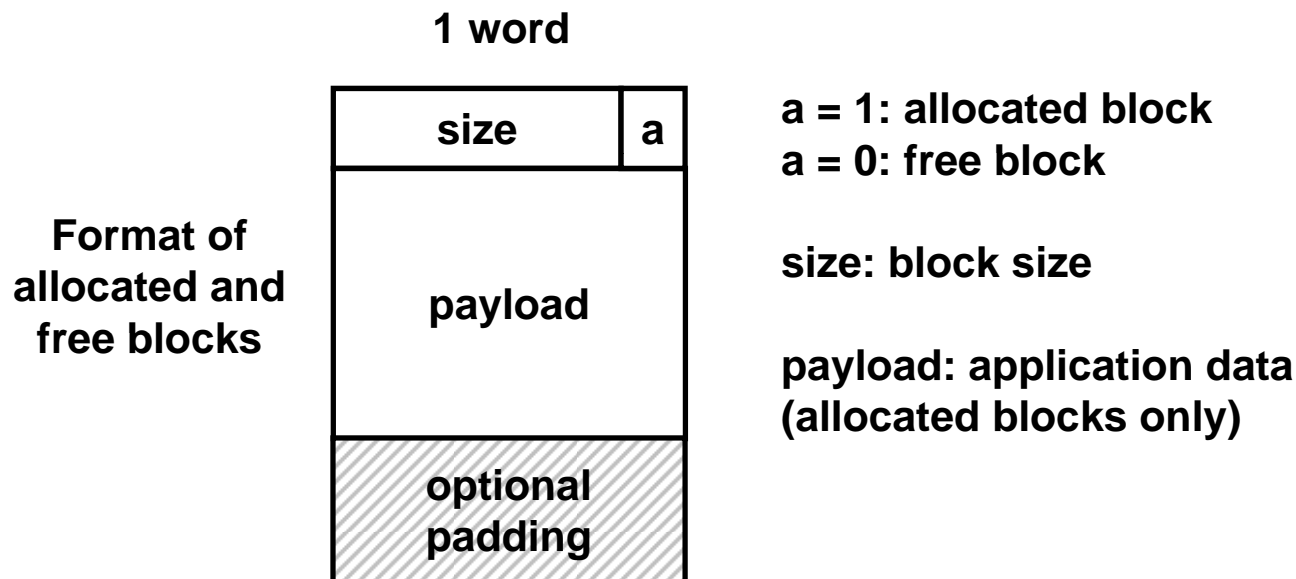
- Different free lists for different size classes

- Method 4: Blocks sorted by size

- Can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key

## 방법 1: 간접 리스트 방식 (Implicit List)

- Need to identify whether each block is free or allocated
  - Can use extra bit
  - Bit can be put in the same word as the size if block sizes are always multiples of two (mask out low order bit when reading size).



## 간접 리스트 : Free 블록 찾기

### ■ *First fit:*

- Search list from beginning, choose first free block that fits

```
p = start;
while ((p < end) &&          \\ not passed end
      ((*p & 1) ||          \\ already allocated
      (*p <= len)))         \\ too small
    p = p + (*p & -2);      \\ goto next block
```

- Can take linear time in total number of blocks (allocated and free)
- In practice it can cause “splinters” at beginning of list

### ■ *Next fit: (Donald Knuth 에 의해 제안됨)*

- Like first-fit, but search list from location of end of previous search
- Research suggests that fragmentation is worse

### ■ *Best fit:*

- Search the list, choose the free block with the closest size that fits
- Keeps fragments small --- usually helps fragmentation
- Will typically run slower than first-fit

### ■ 세가지 검색 방법을 비교할 수 있는가?

## 비트 필드 용법

- How to represent the Header:

- Masks and bitwise operators

```
#define SIZEMASK                (~0x7)
#define PACK(size, alloc)      ((size) | (alloc))
#define getSize(x)              ((x)->size & SIZEMASK)
```


- bitfields

```
struct
{
    unsigned allocated:1;
    unsigned size:31;
} Header;
```

## 간접 리스트 : Free 블록 할당


### ■ Allocating in a free block - *splitting*

- Since allocated space might be smaller than free space, we might want to split the block



```
void addblock(ptr p, int len) {
    int newsize = ((len + 1) >> 1) << 1; // add 1 and round up
    int oldsize = *p & -2;                // mask out low bit
    *p = newsize | 1;                      // set new length
    if (newsize < oldsize)
        *(p+newsize) = oldsize - newsize; // set length in remaining
                                           // part of block
}
```

addblock(p, 4)






## 간접 리스트 : 블록 free 시키기

### ■ Simplest implementation:

#### ● Only need to clear allocated flag

```
void free_block(ptr p) { *p = *p & -2 }
```

#### ● But can lead to "false fragmentation"




*There is enough free space, but the allocator won't be able to find it*

## 간접 리스트 : 연결(Coalescing)

- 다음 또는 이전 블록이 free 하면 함께 연결해서 더 큰 free 블록을 만든다

### Coalescing with next block

```
void free_block(ptr p) {
    *p = *p & -2;           // clear allocated flag
    next = p + *p;           // find next block
    if ((*next & 1) == 0)
        *p = *p + *next;    // add to this block if
                             // not allocated
}
```



### But how do we coalesce with previous block?

## 간접 리스트 : 양방향 연결

### ■ *Boundary tags* [Knuth73]

- Replicate size/allocated word at bottom of free blocks
- Allows us to traverse the "list" backwards, but requires extra space
- Important and general technique!

